

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from scipy.stats import skew
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.svm import SVR
```

```
In [2]: def print_title(title):
        print(f'\n{ '-'*60 }\n\033[1m{title}\033[0m')
def print_section(title):
    print(f'{ '-'*60 }\n{title}\n{ '-'*60 }')
```

(1) Loading and Preprocessing

loading the California Housing dataset using the `fetch_California_Housing` function from `sklearn`

```
In [4]: from sklearn.datasets import fetch_california_housing
data = fetch_california_housing()
```

```
In [5]: x = data.data
y = data.target
feature = data.feature_names
```

```
In [6]: print_section('This data from sklearn dataset contains the average house value as target')
```

```
-----
This data from sklearn dataset contains the average house value as target variable
and the following independent variables (features): average income,
housing average age, average rooms, average bedrooms, population,
average occupation, latitude, and longitude in that order.
-----
```

Converting the dataset into a pandas DataFrame for easier handling.

```
In [8]: ch = pd.DataFrame(x, columns = feature)
ch['MedHouseVal'] = y
```

```
In [9]: #data frame basic exploratory data analysis
print_title("DataFrame Info")
print_section(ch.info())
print_title("DataFrame null values sum")
print_section(ch.isnull().sum())
print_title("DataFrame Description")
print_section(ch.describe())
```

```
-----
DataFrame Info
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
```

```
Data columns (total 9 columns):
#      Column      Non-Null Count  Dtype
---  -
0      MedInc      20640 non-null   float64
1      HouseAge     20640 non-null   float64
2      AveRooms      20640 non-null   float64
3      AveBedrms     20640 non-null   float64
4      Population    20640 non-null   float64
5      AveOccup       20640 non-null   float64
6      Latitude      20640 non-null   float64
7      Longitude      20640 non-null   float64
8      MedHouseVal   20640 non-null   float64
```

```
dtypes: float64(9)
```

```
memory usage: 1.4 MB
```

```
None
```

DataFrame null values sum

```
MedInc      0
HouseAge    0
AveRooms     0
AveBedrms    0
Population   0
AveOccup     0
Latitude     0
Longitude    0
MedHouseVal  0
```

```
dtype: int64
```

DataFrame Description

	MedInc	HouseAge	AveRooms	AveBedrms	Population \
count	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000
mean	3.870671	28.639486	5.429000	1.096675	1425.476744
std	1.899822	12.585558	2.474173	0.473911	1132.462122
min	0.499900	1.000000	0.846154	0.333333	3.000000
25%	2.563400	18.000000	4.440716	1.006079	787.000000
50%	3.534800	29.000000	5.229129	1.048780	1166.000000
75%	4.743250	37.000000	6.052381	1.099526	1725.000000
max	15.000100	52.000000	141.909091	34.066667	35682.000000

	AveOccup	Latitude	Longitude	MedHouseVal
count	20640.000000	20640.000000	20640.000000	20640.000000
mean	3.070655	35.631861	-119.569704	2.068558
std	10.386050	2.135952	2.003532	1.153956
min	0.692308	32.540000	-124.350000	0.149990
25%	2.429741	33.930000	-121.800000	1.196000
50%	2.818116	34.260000	-118.490000	1.797000
75%	3.282261	37.710000	-118.010000	2.647250
max	1243.333333	41.950000	-114.310000	5.000010

```
In [10]: #duplicate row removing
print_section(f"Total Row Befor removing duplicated row: \033[1m{len(ch)}\033[0m")
ch = ch.drop_duplicates()
print_section(f"Total Row After removing duplicated row: \033[1m{len(ch)}\033[0m\nSide B
```

```
Total Row Befor removing duplicated row: 20640
```

Total Row After removing duplicated row: **20640**
Since Both values are same the duplicates rows not available

```
In [11]: # mean and standerd deviation of given dataframe
ch_mean = ch.mean(axis = 0,numeric_only=True)
print_section("\033[1mMean of DataFrame\033[0m")
print_section(ch_mean)
ch_mean = ch.std(axis = 0,numeric_only=True)
print_section("\033[1mStanderd deviation of DataFrame\033[0m")
print_section(ch_mean)
```

Mean of DataFrame



MedInc 3.870671
HouseAge 28.639486
AveRooms 5.429000
AveBedrms 1.096675
Population 1425.476744
AveOccup 3.070655
Latitude 35.631861
Longitude -119.569704
MedHouseVal 2.068558
dtype: float64


Standerd deviation of DataFrame



MedInc 1.899822
HouseAge 12.585558
AveRooms 2.474173
AveBedrms 0.473911
Population 1132.462122
AveOccup 10.386050
Latitude 2.135952
Longitude 2.003532
MedHouseVal 1.153956
dtype: float64



```
In [12]: ### Creating new datafame for outlier visualisation
ch_col = ch.select_dtypes(include='number')
print(ch_col.info())
### syntex to visualise the df to detect outliers for each columns
for i in ch_col.columns:
    sns.boxplot(data=ch,x=i)
    plt.show()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   MedInc          20640 non-null  float64
1   HouseAge        20640 non-null  float64
2   AveRooms        20640 non-null  float64
3   AveBedrms       20640 non-null  float64
4   Population      20640 non-null  float64
5   AveOccup        20640 non-null  float64
6   Latitude        20640 non-null  float64
7   Longitude       20640 non-null  float64
8   MedHouseVal     20640 non-null  float64
dtypes: float64(9)
```











```
In [13]: # creating custom definition to remove outliers using IQR method
def outliers(df):
    for col in df.select_dtypes(include=['int64', 'float64']).columns:
        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1

        lower = Q1 - (1.5*IQR)
        upper = Q3 + (1.5*IQR)

        # Capping
        df[col] = df[col].apply(lambda x: lower if x < lower else upper if x > upper else x)

    return df
```



```
In [14]: #calling custom definition to remove outliers of dataframe 'ch'
new_ch = outliers(ch)
new_ch.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   MedInc          20640 non-null  float64
1   HouseAge        20640 non-null  float64
2   AveRooms        20640 non-null  float64
3   AveBedrms       20640 non-null  float64
4   Population      20640 non-null  float64
5   AveOccup        20640 non-null  float64
6   Latitude        20640 non-null  float64
7   Longitude       20640 non-null  float64
8   MedHouseVal     20640 non-null  float64
dtypes: float64(9)
memory usage: 1.4 MB
```



```
In [15]: ### syntax to visualise the df to detect outliers for each columns after removing outliers
```






```
for i in new_ch.columns:  
    sns.boxplot(data=new_ch,x=i)  
    plt.show()
```











```
In [16]: #minmax scaling
numerical_col = new_ch.select_dtypes(include=['number'])
scl_Minmax = MinMaxScaler()
scl_dta_mimax = scl_Minmax.fit_transform(numerical_col)
ch_scl_minmax = pd.DataFrame(scl_dta_mimax, columns=numerical_col.columns, index=new_ch.index)
New_ch_minmax = pd.concat([new_ch.drop(columns=numerical_col.columns), ch_scl_minmax], axis=1)
print_title("New df after MinMax Scaling")
print_section(f'\n{New_ch_minmax}')
```

New df after MinMax Scaling

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	\
0	1.000000	0.784314	0.769532	0.422435	0.101950	0.411895	
1	1.000000	0.392157	0.653814	0.283508	0.766379	0.281190	
2	0.899426	1.000000	0.971808	0.555229	0.157558	0.484240	
3	0.684562	1.000000	0.588542	0.554193	0.177373	0.409663	
4	0.445394	1.000000	0.660596	0.575654	0.179610	0.302194	
...	
20635	0.141140	0.470588	0.468806	0.715445	0.269096	0.413376	
20636	0.273774	0.333333	0.634564	1.000000	0.112816	0.578240	
20637	0.159734	0.313725	0.493639	0.680021	0.320869	0.344471	
20638	0.181988	0.333333	0.512869	0.818676	0.235858	0.285110	
20639	0.251387	0.294118	0.501267	0.792844	0.442314	0.429908	

	Latitude	Longitude	MedHouseVal
0	0.567481	0.211155	0.936218
1	0.565356	0.212151	0.734897
2	0.564293	0.210159	0.721205
3	0.564293	0.209163	0.698099
4	0.564293	0.209163	0.700025
...
20635	0.737513	0.324701	0.135000
20636	0.738576	0.312749	0.132861
20637	0.732200	0.311753	0.165380
20638	0.732200	0.301793	0.149121
20639	0.725824	0.309761	0.159176

[20640 rows x 9 columns]

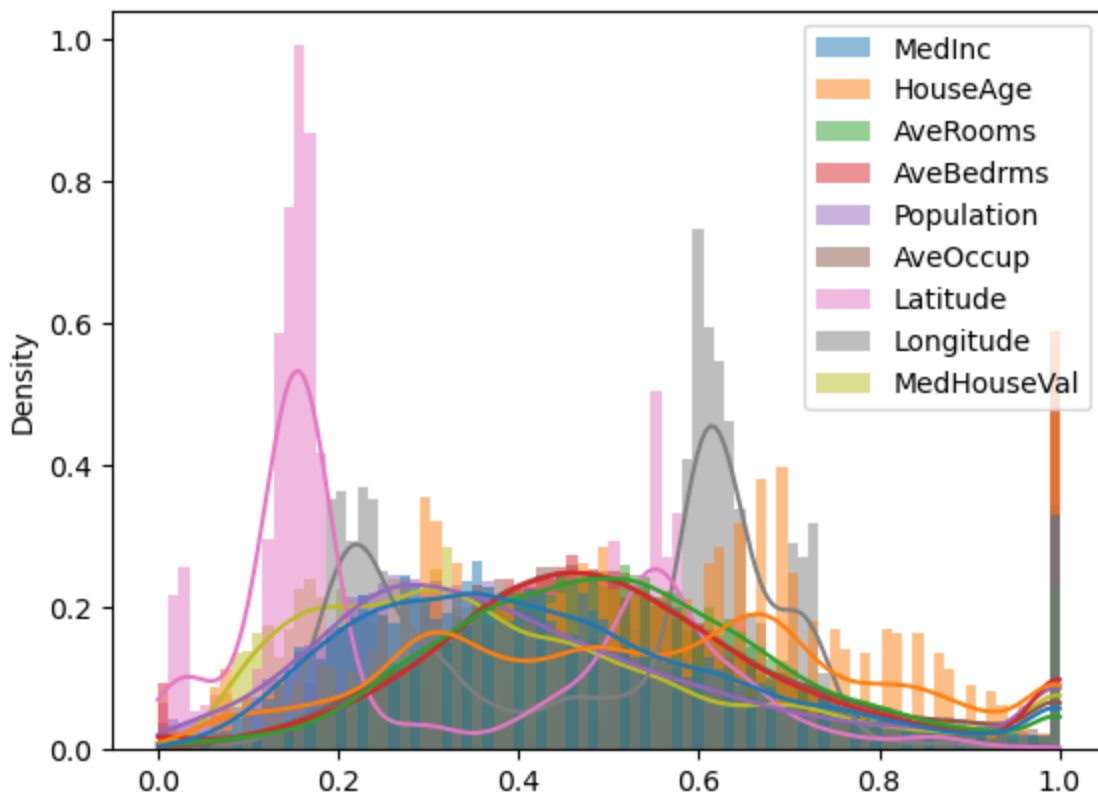
```
In [17]: skewness = New_ch_minmax.skew()
print_section("\033[1mSkewness of Features in dataframe after Scaling\033[0m")
print_section(skewness)
```

Skewness of Features in dataframe after Scaling

```
MedInc          0.735618
HouseAge        0.060331
AveRooms        0.348567
AveBedrms       0.462645
Population      0.842247
AveOccup        0.510453
Latitude        0.465953
Longitude       -0.297801
MedHouseVal     0.912330
dtype: float64
```


```
In [18]: sns.histplot(New_ch_minmax,kde=True,stat='density',linewidth=0)
print_section("\033[1mHistplot of dataframe after Scaling\033[0m")
plt.show()
```

Histplot of dataframe after Scaling



```
In [19]: corr_ch = New_ch_minmax.corr()
```

```
In [20]: sns.heatmap(corr_ch,annot=True)
plt.show()
skewness = New_ch_minmax.skew()
print_section("\033[1mSkewness of Features in dataframe after Scaling\033[0m")
print_section(skewness)
```



Skewness of Features in dataframe after Scaling

```

MedInc          0.735618
HouseAge        0.060331
AveRooms        0.348567
AveBedrms       0.462645
Population      0.842247
AveOccup        0.510453
Latitude        0.465953
Longitude       -0.297801
MedHouseVal     0.912330
dtype: float64

```

The Histplot of after scaling and before scaling gives a clear visual of above steps of data preprocessing is necessary for the dataset and Highly skewed features have good co-relation with the target variable 'MedHouseVal' and less co-related features have good skewness. Therefore, the dataset 'New_ch_minmax' can be concluded to be fit for the next process

(2) Regression Algorithm Implementation

```

In [23]: Cal_H=New_ch_minmax #renaming for further Analysis
# Feature selection
X=Cal_H[['MedInc','HouseAge','AveRooms','AveBedrms','Population','AveOccup','Latitude','Longitude']]
Y=Cal_H['MedHouseVal']

```

```

In [24]: #dependent variable
Y

```

```

0          0.936218

```

```
Out[24]: 1      0.734897
          2      0.721205
          3      0.698099
          4      0.700025
          ...
        20635    0.135000
        20636    0.132861
        20637    0.165380
        20638    0.149121
        20639    0.159176
Name: MedHouseVal, Length: 20640, dtype: float64
```

Linear Regression

```
In [26]: #syntax for training and testing for linear regression
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=42)
```

```
In [27]: #syntax for linear regression modeling
model = LinearRegression()
model.fit(X_train, y_train)
```

```
Out[27]: ▼ LinearRegression ⓘ ?
LinearRegression()
```

```
In [28]: #syntax for model prediction
y_pred = model.predict(X_test)
print(y_pred)

[0.02279585 0.38853306 0.69908632 ... 0.83300977 0.23446578 0.32665012]
```

```
In [29]: #syntax for model score
Score_LinearRegression = model.score(X_test, y_test)
```

Decision Tree Regressor

```
In [31]: #syntax for training and testing for Decision Tree Regressor
X1_train, X1_test, y1_train, y1_test = train_test_split(X, Y, test_size=0.2, random_state=42)
```

```
In [32]: #syntax for Decision Tree Regressor modeling
model1 = DecisionTreeRegressor(random_state=42)
model1.fit(X1_train, y1_train)
```

```
Out[32]: ▼ DecisionTreeRegressor ⓘ ?
DecisionTreeRegressor(random_state=42)
```

```
In [33]: #syntax for model prediction
y1_pred = model1.predict(X1_test)
print(y1_pred)

[0.08729102 0.22528446 1.          ... 1.          0.10911324 0.27042651]
```

```
In [34]: #syntax for model score
Score_DecisionTreeRegressor = model1.score(X1_test, y1_test)
```

Random Forest Regressor

```
In [36]: #syntax for training and testing for Random Forest Regressor
```



```
X2_train, X2_test, y2_train, y2_test = train_test_split(X, Y, test_size=0.2, random_stat
```

```
In [37]: #syntes for Random Forest Regressor modeling
model2 = RandomForestRegressor(random_state=42)
model2.fit(X2_train, y2_train)
```

```
Out[37]: ▼ RandomForestRegressor ⓘ ?
RandomForestRegressor(random_state=42)
```

```
In [38]: #syntex for model prediction
y2_pred = model2.predict(X2_test)
print(y2_pred)

[0.0779203  0.12537293 0.98714789 ... 0.92966881 0.12087156 0.32966741]
```

```
In [39]: #sytex for model score
Score_RandomForestRegressor = model2.score(X2_test,y2_test)
```

Gradient Boosting Regressor

```
In [41]: #syntex for traning and testing for Gradient Boosting Regressor
X3_train, X3_test, y3_train, y3_test = train_test_split(X, Y, test_size=0.2, random_stat
```

```
In [42]: #syntes for Gradient Boosting Regressor modeling
model3 = GradientBoostingRegressor(random_state=42)
model3.fit(X3_train, y3_train)
```

```
Out[42]: ▼ GradientBoostingRegressor ⓘ ?
GradientBoostingRegressor(random_state=42)
```

```
In [43]: #syntex for model prediction
y3_pred = model3.predict(X3_test)
print(y2_pred)

[0.0779203  0.12537293 0.98714789 ... 0.92966881 0.12087156 0.32966741]
```

```
In [115... #sytex for model score
Score_GradientBootingRegressor = model3.score(X3_test,y3_test)
```

Support Vector Regressor (SVR)

```
In [46]: scaler_x = StandardScaler()
x_scaled = scaler_x.fit_transform(X)
scaler_y = StandardScaler()
y_scaled = scaler_y.fit_transform(Y.values.reshape(-1, 1))
```

```
In [47]: #syntex for traning and testing for Support Vector Regressor
X4_train, X4_test, y4_train, y4_test = train_test_split(x_scaled, y_scaled, test_size=0.
```

```
In [48]: #syntes for Support Vector Regressor modeling
model4 = SVR(kernel='rbf', C=1.0, epsilon=0.1)
model4.fit(X4_train, y4_train.ravel())
```

```
Out[48]: ▼ SVR ⓘ ?
SVR()
```

```
In [49]: #syntax for model prediction
y4_pred = model4.predict(X4_test)
print(y4_pred)

[-1.40585189 -0.53264762  2.2693782   ...   2.14964491 -1.29540304
 -0.27994881]
```

```
In [50]: #syntax for model score
Score_SVR = model4.score(X4_test,y4_test)
```

```
In [121]: print(f'Accuracy of Linear Regression is \033[1m{int(100*(Score_LinearRegression))}%\033
print(f'The RandomForest Regression Accuracy rate is \033[1m{int(100*(Score_RandomForest

Accuracy of Linear Regression is 65%
Accuracy of DecisionTree Regression is 62%
Accuracy of RandomForest Regression is 80%
Accuracy of GradientBoosting Regression is 77%
and Accuracy of SVR Regression is 76%
The RandomForest Regression Accuracy rate is 80%, it is higher than other methond
```

(3) Model Evelution Comparison

Evaluateing the performance of Mean Squared Error(MSE), Mean Absolute error(MAE) and R-Squared Score

Evaluating Linear Regression model using following metrics

```
In [55]: mae_LR = mean_absolute_error(y_test,y_pred)
mse_LR = mean_squared_error(y_test,y_pred)
r2_LR = r2_score(y_test,y_pred)
print_title(f'The Evaluated performace of MSE, MAE and R-Squred Score\nfor Linear Regres
print_section(f'Mean Squared Error: {mse_LR}\nMean Absolute Error: {mae_LR}\nR-Squared S

-----
The Evaluated performace of MSE, MAE and R-Squred Score
for Linear Regression model
-----

Mean Squared Error: 0.020249838675798007
Mean Absolute Error: 0.10583073170383961
R-Squared Score: 0.6500913755505537
-----
```

Evaluating Decidion Tree Regressor model using following metrics

```
In [57]: mae_DTR = mean_absolute_error(y1_test,y1_pred)
mse_DTR = mean_squared_error(y1_test,y1_pred)
r2_DTR = r2_score(y1_test,y1_pred)
print_title(f'The Evaluated performace of MSE, MAE and R-Squred Score\nfor Decidion Tree
print_section(f'Mean Squared Error: {mse_DTR}\nMean Absolute Error: {mae_DTR}\nR-Squared

-----
The Evaluated performace of MSE, MAE and R-Squred Score
for Decidion Tree Regresso model
-----

Mean Squared Error: 0.02195135377872779
Mean Absolute Error: 0.09570921090834345
R-Squared Score: 0.6206899161770678
-----
```

Evaluating Random Forest Regressor model using following metrics

```
In [59]: mae_RFR = mean_absolute_error(y2_test,y2_pred)
mse_RFR = mean_squared_error(y2_test,y2_pred)
r2_RFR = r2_score(y2_test,y2_pred)
print_title(f'The Evaluated performace of MSE, MAE and R-Squared Score\nfor Random Forest
print_section(f'Mean Squared Error: {mse_RFR}\nMean Absolute Error: {mae_RFR}\nR-Squared
```

```
-----
The Evaluated performace of MSE, MAE and R-Squared Score
for Random Forest Regressor model
-----
```

```
Mean Squared Error: 0.01114867521661751
Mean Absolute Error: 0.06943318886675323
R-Squared Score: 0.8073556203614282
-----
```

Evaluating Gradient Boosting Regressor model using following metrics

```
In [61]: mae_GBR = mean_absolute_error(y3_test,y3_pred)
mse_GBR = mean_squared_error(y3_test,y3_pred)
r2_GBR = r2_score(y3_test,y3_pred)
print_title(f'The Evaluated performace of MSE, MAE and R-Squared Score\nfor Gradient Boos
print_section(f'Mean Squared Error: {mse_GBR}\nMean Absolute Error: {mae_GBR}\nR-Squared
```

```
-----
The Evaluated performace of MSE, MAE and R-Squared Score
for Gradient Boosting Regressor model
-----
```

```
Mean Squared Error: 0.012799214573888404
Mean Absolute Error: 0.07796846596239415
R-Squared Score: 0.7788350002543363
-----
```

Evaluating Support Vector Regressor model using following metrics

```
In [63]: mae_SVR = mean_absolute_error(y4_test,y4_pred)
mse_SVR = mean_squared_error(y4_test,y4_pred)
r2_SVR = r2_score(y4_test,y4_pred)
print_title(f'The Evaluated performace of MSE, MAE and R-Squared Score\nfor Support Vecto
print_section(f'Mean Squared Error: {mse_SVR}\nMean Absolute Error: {mae_SVR}\nR-Squared
```

```
-----
The Evaluated performace of MSE, MAE and R-Squared Score
for Support Vector Regressor model
-----
```

```
Mean Squared Error: 0.23314051024564758
Mean Absolute Error: 0.3249791281149936
R-Squared Score: 0.7636512832262522
-----
```

```
In [ ]:
```

Comparing and identifying the results of all models

```
In [65]: R_score = {'model':['Linear Regression','Decidion Tree Regressor','Random Forest Regress
          'Gradient Boosting Regressor','Support Vector Regressor'],
          'R2_Score %':[100*r2_LR,100*r2_DTR,100*r2_RFR,100*r2_GBR,100*r2_SVR]}
R_score = pd.DataFrame(R_score)
highh= R_score['R2_Score %'].max()
low= R_score['R2_Score %'].min()
Iden_model = R_score.loc[R_score['R2_Score %'] == highh, 'model'].item()
wrost_model = R_score.loc[R_score['R2_Score %'] == low, 'model'].item()
print_title('                Result Comparison using R-Squared Score')
print_section(f'Compared and Identified model using R2_Score metric\nof \033[1m{Iden_mod
```

Result Comparison using R-Squared Score

Compared and Identified model using R2_Score metric
of **Random Forest Regressor** its R-Squared Score
is **80%** and it has best-performing algorithm
with justification. And the Worst model
is **Decision Tree Regressor** its R-Squared score is **62%**