

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from scipy.stats import skew
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.svm import SVR
```

```
In [2]: def print_title(title):
        print(f'\n{ '-'*60 }\n\033[1m{title}\033[0m')
def print_section(title):
    print(f'\n{ '-'*60 }\n{title}\n{ '-'*60 }')
```

## (1) Loading and Preprocessing

loading the California Housing dataset using the `fetch_California_Housing` function from `sklearn`

```
In [4]: from sklearn.datasets import fetch_california_housing
data = fetch_california_housing()
```

```
In [5]: x = data.data
y = data.target
feature = data.feature_names
```

```
In [6]: print_section('This data from sklearn dataset contains the average house value as target')
```

```
-----
This data from sklearn dataset contains the average house value as target variable
and the following independent variables (features): average income,
housing average age, average rooms, average bedrooms, population,
average occupation, latitude, and longitude in that order.
-----
```

Converting the dataset into a pandas DataFrame for easier handling.

```
In [8]: ch = pd.DataFrame(x, columns = feature)
ch['MedHouseVal'] = y
```

```
In [9]: #data frame basic exploratory data analysis
print_title("DataFrame Info")
print_section(ch.info())
print_title("DataFrame null values sum")
print_section(ch.isnull().sum())
print_title("DataFrame Description")
print_section(ch.describe())
```

```
-----
DataFrame Info
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
```

```
Data columns (total 9 columns):
#      Column      Non-Null Count  Dtype
---  -
0      MedInc      20640 non-null   float64
1      HouseAge     20640 non-null   float64
2      AveRooms      20640 non-null   float64
3      AveBedrms     20640 non-null   float64
4      Population    20640 non-null   float64
5      AveOccup       20640 non-null   float64
6      Latitude       20640 non-null   float64
7      Longitude      20640 non-null   float64
8      MedHouseVal    20640 non-null   float64
```

```
dtypes: float64(9)
```

```
memory usage: 1.4 MB
```

```
None
```

#### DataFrame null values sum

```
MedInc      0
HouseAge    0
AveRooms     0
AveBedrms    0
Population   0
AveOccup     0
Latitude     0
Longitude    0
MedHouseVal  0
```

```
dtype: int64
```

#### DataFrame Description

|       | MedInc       | HouseAge     | AveRooms     | AveBedrms    | Population \ |
|-------|--------------|--------------|--------------|--------------|--------------|
| count | 20640.000000 | 20640.000000 | 20640.000000 | 20640.000000 | 20640.000000 |
| mean  | 3.870671     | 28.639486    | 5.429000     | 1.096675     | 1425.476744  |
| std   | 1.899822     | 12.585558    | 2.474173     | 0.473911     | 1132.462122  |
| min   | 0.499900     | 1.000000     | 0.846154     | 0.333333     | 3.000000     |
| 25%   | 2.563400     | 18.000000    | 4.440716     | 1.006079     | 787.000000   |
| 50%   | 3.534800     | 29.000000    | 5.229129     | 1.048780     | 1166.000000  |
| 75%   | 4.743250     | 37.000000    | 6.052381     | 1.099526     | 1725.000000  |
| max   | 15.000100    | 52.000000    | 141.909091   | 34.066667    | 35682.000000 |

|       | AveOccup     | Latitude     | Longitude    | MedHouseVal  |
|-------|--------------|--------------|--------------|--------------|
| count | 20640.000000 | 20640.000000 | 20640.000000 | 20640.000000 |
| mean  | 3.070655     | 35.631861    | -119.569704  | 2.068558     |
| std   | 10.386050    | 2.135952     | 2.003532     | 1.153956     |
| min   | 0.692308     | 32.540000    | -124.350000  | 0.149990     |
| 25%   | 2.429741     | 33.930000    | -121.800000  | 1.196000     |
| 50%   | 2.818116     | 34.260000    | -118.490000  | 1.797000     |
| 75%   | 3.282261     | 37.710000    | -118.010000  | 2.647250     |
| max   | 1243.333333  | 41.950000    | -114.310000  | 5.000010     |

```
In [10]: #duplicate row removing
print_section(f"Total Row Befor removing duplicated row: \033[1m{len(ch)}\033[0m")
ch = ch.drop_duplicates()
print_section(f"Total Row After removing duplicated row: \033[1m{len(ch)}\033[0m\nSize B
```

```
Total Row Befor removing duplicated row: 20640
```

Total Row After removing duplicated row: **20640**  
Since Both values are same the duplicates rows not available  
-----

```
In [11]: # mean and standerd deviation of given dataframe
ch_mean = ch.mean(axis = 0,numeric_only=True)
print_section("\033[1mMean of DataFrame\033[0m")
print_section(ch_mean)
ch_mean = ch.std(axis = 0,numeric_only=True)
print_section("\033[1mStanderd deviation of DataFrame\033[0m")
print_section(ch_mean)
```

-----  
**Mean of DataFrame**  
-----

-----  
MedInc 3.870671  
HouseAge 28.639486  
AveRooms 5.429000  
AveBedrms 1.096675  
Population 1425.476744  
AveOccup 3.070655  
Latitude 35.631861  
Longitude -119.569704  
MedHouseVal 2.068558  
dtype: float64  
-----



-----  
**Standerd deviation of DataFrame**  
-----



-----  
MedInc 1.899822  
HouseAge 12.585558  
AveRooms 2.474173  
AveBedrms 0.473911  
Population 1132.462122  
AveOccup 10.386050  
Latitude 2.135952  
Longitude 2.003532  
MedHouseVal 1.153956  
dtype: float64  
-----



```
In [12]: ### Creating new datafame for outlier visualisation
ch_col = ch.select_dtypes(include='number')
print(ch_col.info())
### syntex to visualise the df to detect outliers for each columns
for i in ch_col.columns:
    sns.boxplot(data=ch,x=i)
    plt.show()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   MedInc      20640 non-null  float64
1   HouseAge    20640 non-null  float64
2   AveRooms    20640 non-null  float64
3   AveBedrms   20640 non-null  float64
4   Population  20640 non-null  float64
5   AveOccup    20640 non-null  float64
6   Latitude    20640 non-null  float64
7   Longitude   20640 non-null  float64
8   MedHouseVal 20640 non-null  float64
dtypes: float64(9)
```


memory usage: 1.4 MB  
None











```
In [13]: # creating custom definition to remove outliers using IQR method
def outliers(df):
    for col in df.select_dtypes(include=['int64', 'float64']).columns:
        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1

        lower = Q1 - (1.5*IQR)
        upper = Q3 + (1.5*IQR)

        # Capping
        df[col] = df[col].apply(lambda x: lower if x < lower else upper if x > upper else x)

    return df
```



```
In [14]: #calling custom definition to remove outliers of dataframe 'ch'
new_ch = outliers(ch)
new_ch.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   MedInc          20640 non-null  float64
1   HouseAge        20640 non-null  float64
2   AveRooms        20640 non-null  float64
3   AveBedrms       20640 non-null  float64
4   Population      20640 non-null  float64
5   AveOccup        20640 non-null  float64
6   Latitude        20640 non-null  float64
7   Longitude       20640 non-null  float64
8   MedHouseVal     20640 non-null  float64
dtypes: float64(9)
memory usage: 1.4 MB
```



```
In [15]: ### syntax to visualise the df to detect outliers for each columns after removing outliers
```





```
for i in new_ch.columns:  
    sns.boxplot(data=new_ch,x=i)  
    plt.show()
```











[20640 rows x 9 columns]


```
In [17]: skewness = New_ch_minmax.skew()
print_section("\033[1mSkewness of Features in dataframe after Scaling\033[0m")
print_section(skewness)
```

#### Skewness of Features in dataframe after Scaling

```
MedInc          0.735618
HouseAge        0.060331
AveRooms        0.348567
AveBedrms       0.462645
Population      0.842247
AveOccup        0.510453
Latitude        0.465953
Longitude       -0.297801
MedHouseVal     0.912330
dtype: float64
```


```
In [18]: sns.histplot(New_ch_minmax,kde=True,stat='density',linewidth=0)
print_section("\033[1mHistplot of dataframe after Scaling\033[0m")
plt.show()
```

#### Histplot of dataframe after Scaling



```
In [19]: corr_ch = New_ch_minmax.corr()
```

```
In [20]: sns.heatmap(corr_ch,annot=True)
plt.show()
skewness = New_ch_minmax.skew()
print_section("\033[1mSkewness of Features in dataframe after Scaling\033[0m")
print_section(skewness)
```



#### Skewness of Features in dataframe after Scaling

```

MedInc      0.735618
HouseAge    0.060331
AveRooms    0.348567
AveBedrms   0.462645
Population  0.842247
AveOccup    0.510453
Latitude    0.465953
Longitude   -0.297801
MedHouseVal 0.912330
dtype: float64

```

The Histplot of after scaling and before scaling gives a clear visual of above steps of data preprocessing is necessary for the dataset and Highly skewed features have good co-relation with the target variable 'MedHouseVal' and less co-related features have good skewness. Therefore, the dataset 'New\_ch\_minmax' can be concluded to be fit for the next process

## (2) Regression Algorithm Implementation

```

In [23]: Cal_H = New_ch_minmax #renaming for further Analysis
# Feature selection
X = Cal_H[['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup', 'Latitude', 'Longitude']]
Y = Cal_H['MedHouseVal']

```

```

In [24]: #dependent variable
Y

```

```

0      0.936218

```

```
Out[24]: 1      0.734897
         2      0.721205
         3      0.698099
         4      0.700025
         ...
        20635    0.135000
        20636    0.132861
        20637    0.165380
        20638    0.149121
        20639    0.159176
Name: MedHouseVal, Length: 20640, dtype: float64
```

## Linear Regression

```
In [26]: #syntax for training and testing for linear regression
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=42)
```

```
In [27]: #syntax for linear regression modeling
model = LinearRegression()
model.fit(X_train, y_train)
```

```
Out[27]: ▼ LinearRegression ⓘ ?
LinearRegression()
```

```
In [28]: #syntax for model prediction
y_pred = model.predict(X_test)
print(y_pred)

[0.02279585 0.38853306 0.69908632 ... 0.83300977 0.23446578 0.32665012]
```

```
In [29]: #syntax for model score
Score_LinearRegression = model.score(X_test, y_test)
```

## Decision Tree Regressor

```
In [31]: #syntax for training and testing for Decision Tree Regressor
X1_train, X1_test, y1_train, y1_test = train_test_split(X, Y, test_size=0.2, random_state=42)
```

```
In [32]: #syntax for Decision Tree Regressor modeling
model1 = DecisionTreeRegressor(random_state=42)
model1.fit(X1_train, y1_train)
```

```
Out[32]: ▼ DecisionTreeRegressor ⓘ ?
DecisionTreeRegressor(random_state=42)
```

```
In [33]: #syntax for model prediction
y1_pred = model1.predict(X1_test)
print(y1_pred)

[0.08729102 0.22528446 1.          ... 1.          0.10911324 0.27042651]
```

```
In [34]: #syntax for model score
Score_DecisionTreeRegressor = model1.score(X1_test, y1_test)
```

## Random Forest Regressor

```
In [36]: #syntax for training and testing for Random Forest Regressor
```



```
X2_train, X2_test, y2_train, y2_test = train_test_split(X, Y, test_size=0.2, random_stat
```

```
In [37]: #syntes for Random Forest Regressor modeling
model2 = RandomForestRegressor(random_state=42)
model2.fit(X2_train, y2_train)
```

```
Out[37]: ▼      RandomForestRegressor ⓘ ⓘ
RandomForestRegressor(random_state=42)
```

```
In [38]: #syntex for model prediction
y2_pred = model2.predict(X2_test)
print(y2_pred)

[0.0779203  0.12537293 0.98714789 ... 0.92966881 0.12087156 0.32966741]
```

```
In [39]: #sytex for model score
Score_RandomForestRegressor = model1.score(X2_test,y2_test)
```

## Gradient Boosting Regressor

```
In [41]: #syntex for traning and testing for Gradient Boosting Regressor
X3_train, X3_test, y3_train, y3_test = train_test_split(X, Y, test_size=0.2, random_stat
```

```
In [42]: #syntes for Gradient Boosting Regressor modeling
model3 = GradientBoostingRegressor(random_state=42)
model3.fit(X3_train, y3_train)
```

```
Out[42]: ▼      GradientBoostingRegressor ⓘ ⓘ
GradientBoostingRegressor(random_state=42)
```

```
In [43]: #syntex for model prediction
y3_pred = model3.predict(X3_test)
print(y2_pred)

[0.0779203  0.12537293 0.98714789 ... 0.92966881 0.12087156 0.32966741]
```

```
In [44]: #sytex for model score
Score_GradientBootingRegressor = model1.score(X3_test,y3_test)
```

## Support Vector Regressor (SVR)

```
In [46]: scaler_x = StandardScaler()
x_scaled = scaler_x.fit_transform(X)
scaler_y = StandardScaler()
y_scaled = scaler_y.fit_transform(Y.values.reshape(-1, 1))
```

```
In [47]: #syntex for traning and testing for Support Vector Regressor
X4_train, X4_test, y4_train, y4_test = train_test_split(x_scaled, y_scaled, test_size=0.
```

```
In [48]: #syntes for Support Vector Regressor modeling
model4 = SVR(kernel='rbf', C=1.0, epsilon=0.1)
model4.fit(X4_train, y4_train.ravel())
```

```
Out[48]: ▼      SVR ⓘ ⓘ
SVR()
```

```
In [49]: #syntax for model prediction
y4_pred = model4.predict(X4_test)
print(y4_pred)

[-1.40585189 -0.53264762  2.2693782   ...   2.14964491 -1.29540304
 -0.27994881]

In [50]: #syntax for model score
Score_SVR = model4.score(X4_test,y4_test)

In [51]: print(f'Accuracy of Linear Regression is \033[1m{int(100*(Score_LinearRegression))}%\033
print(f'The SVR Accuracy rate is \033[1m{int(100*(Score_SVR))}%\033[0m, it is higher tha

Accuracy of Linear Regression is 65%
Accuracy of DecisionTree Regression is 62%
Accuracy of RandomForest Regression is 62%
Accuracy of GradientBoosting Regression is 62%
and Accuracy of SVR Regression is 76%
The SVR Accuracy rate is 76%, it is higher than other methond
```

### (3) Model Evelution Comparison

#### Evaluating the performance of Mean Squared Error(MSE), Mean Absolute error(MAE) and R-Squared Score

##### Evaluating Linear Regression model using following metrics

```
In [55]: mae_LR = mean_absolute_error(y_test,y_pred)
mse_LR = mean_squared_error(y_test,y_pred)
r2_LR = r2_score(y_test,y_pred)
print_title(f'The Evaluated performace of MSE, MAE and R-Squared Score\nfor Linear Regres
print_section(f'Mean Squared Error: {mse_LR}\nMean Absolute Error: {mae_LR}\nR-Squared S

-----

The Evaluated performace of MSE, MAE and R-Squared Score
for Linear Regression model

-----

Mean Squared Error: 0.020249838675798007
Mean Absolute Error: 0.10583073170383961
R-Squared Score: 0.6500913755505537
-----
```

##### Evaluating Decidion Tree Regressor model using following metrics

```
In [57]: mae_DTR = mean_absolute_error(y1_test,y1_pred)
mse_DTR = mean_squared_error(y1_test,y1_pred)
r2_DTR = r2_score(y1_test,y1_pred)
print_title(f'The Evaluated performace of MSE, MAE and R-Squared Score\nfor Decidion Tree
print_section(f'Mean Squared Error: {mse_DTR}\nMean Absolute Error: {mae_DTR}\nR-Squared

-----

The Evaluated performace of MSE, MAE and R-Squared Score
for Decidion Tree Regresso model

-----

Mean Squared Error: 0.02195135377872779
Mean Absolute Error: 0.09570921090834345
R-Squared Score: 0.6206899161770678
-----
```

##### Evaluating Random Forest Regressor model using following metrics

```
In [59]: mae_RFR = mean_absolute_error(y2_test,y2_pred)
mse_RFR = mean_squared_error(y2_test,y2_pred)
r2_RFR = r2_score(y2_test,y2_pred)
print_title(f'The Evaluated performace of MSE, MAE and R-Squared Score\nfor Random Forest
print_section(f'Mean Squared Error: {mse_RFR}\nMean Absolute Error: {mae_RFR}\nR-Squared
```

```
-----
The Evaluated performace of MSE, MAE and R-Squared Score
for Random Forest Regressor model
-----
```

```
Mean Squared Error: 0.01114867521661751
Mean Absolute Error: 0.06943318886675323
R-Squared Score: 0.8073556203614282
-----
```

## Evaluating Gradient Boosting Regressor model using following metrics

```
In [61]: mae_GBR = mean_absolute_error(y3_test,y3_pred)
mse_GBR = mean_squared_error(y3_test,y3_pred)
r2_GBR = r2_score(y3_test,y3_pred)
print_title(f'The Evaluated performace of MSE, MAE and R-Squared Score\nfor Gradient Boos
print_section(f'Mean Squared Error: {mse_GBR}\nMean Absolute Error: {mae_GBR}\nR-Squared
```

```
-----
The Evaluated performace of MSE, MAE and R-Squared Score
for Gradient Boosting Regressor model
-----
```

```
Mean Squared Error: 0.012799214573888404
Mean Absolute Error: 0.07796846596239415
R-Squared Score: 0.7788350002543363
-----
```

## Evaluating Support Vector Regressor model using following metrics

```
In [63]: mae_SVR = mean_absolute_error(y4_test,y4_pred)
mse_SVR = mean_squared_error(y4_test,y4_pred)
r2_SVR = r2_score(y4_test,y4_pred)
print_title(f'The Evaluated performace of MSE, MAE and R-Squared Score\nfor Support Vecto
print_section(f'Mean Squared Error: {mse_SVR}\nMean Absolute Error: {mae_SVR}\nR-Squared
```

```
-----
The Evaluated performace of MSE, MAE and R-Squared Score
for Support Vector Regressor model
-----
```

```
Mean Squared Error: 0.23314051024564758
Mean Absolute Error: 0.3249791281149936
R-Squared Score: 0.7636512832262522
-----
```

```
In [ ]:
```

## Comparing and identifying the results of all models

```
In [69]: R_score = {'model':['Linear Regression','Decidion Tree Regressor','Random Forest Regress
          'Gradient Boosting Regressor','Support Vector Regressor'],
          'R2_Score %':[100*r2_LR,100*r2_DTR,100*r2_RFR,100*r2_GBR,100*r2_SVR]}
R_score = pd.DataFrame(R_score)
high= R_score['R2_Score %'].max()
low= R_score['R2_Score %'].min()
Iden_model = R_score.loc[R_score['R2_Score %'] == high, 'model'].item()
wrost_model = R_score.loc[R_score['R2_Score %'] == low, 'model'].item()
print_title('          "Result Comparison using R-Squared Score")
print_section(f'Compared and Identified model using R2_Score metric\nof \033[1m{Iden_mod
```

### Result Comparison using R-Squared Score

Compared and Identified model using R2\_Score metric  
of **Random Forest Regressor** its R-Squared Score  
is **80%** and it has best-performing algorithm  
with justification. And the Worst model  
is **Decision Tree Regressor** its R-Squared score is **62%**